

Chatbot lab

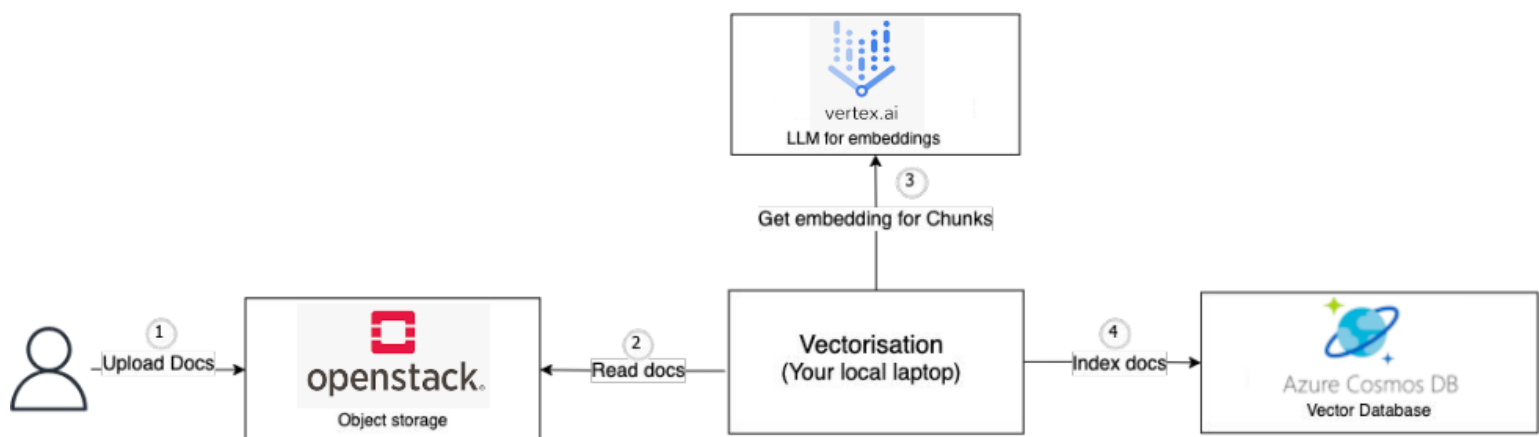
Authors: Olivia Manz and Samuel Roland in **Group D**

The whole project is available on this Github repository: <https://github.com/samuelroland/CloudSys-labs/tree/main/lab3>

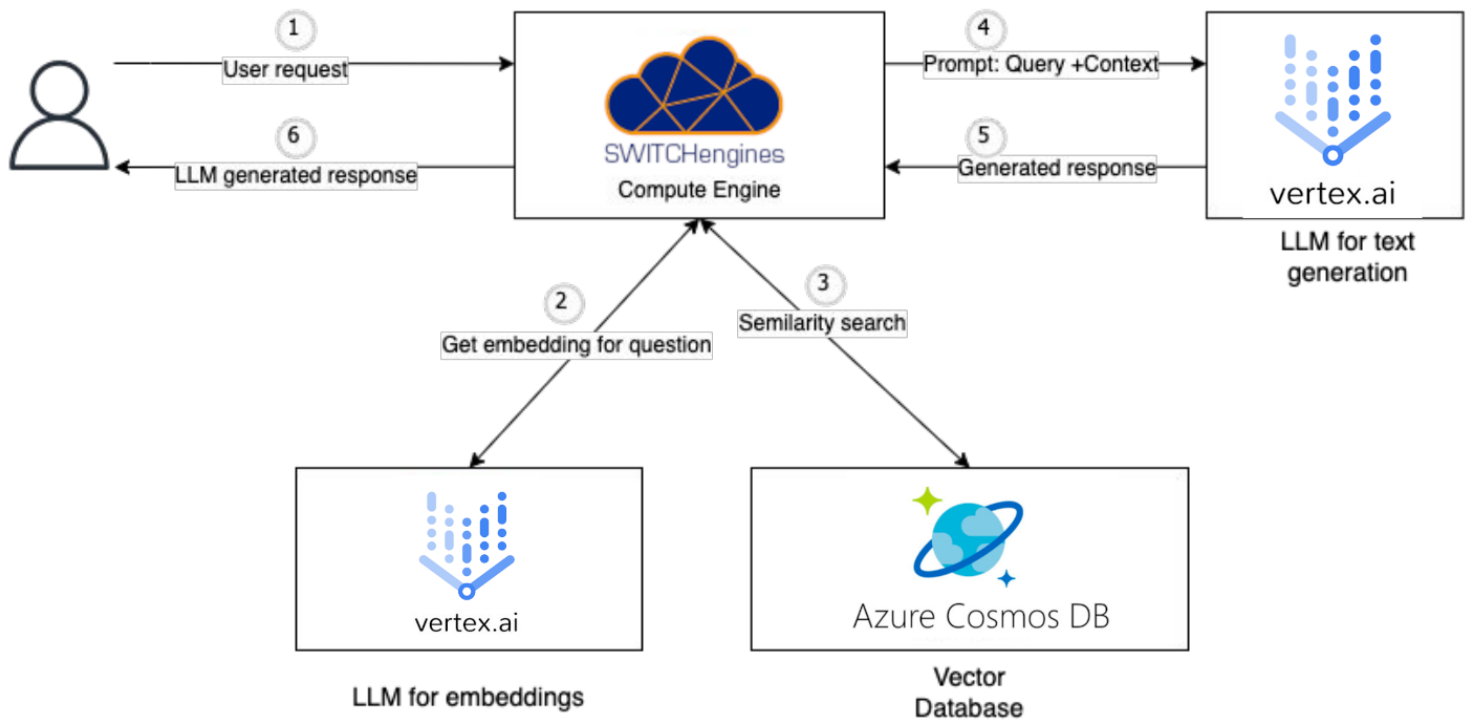
Goal

Using proprietary APIs of Cloud providers from a single online chatbot. Based on the slides schemas, we updated the 2 architecture diagrams.

First step: Setup an S3 bucket on OpenStack, be able to download documents, send them to Google Vertex AI to get their embeddings equivalent. Then, we want to store these floats vectors into an Azure Cosmos DB database.



Second step: Deploy a chatbot based on given code, that has a streamlit chat interface that allows to ask a question about the vectorized documents. First the user request the bot has received, is sent to Google Vertex AI again to a model called `gemini-embedding-001`, the returned embedding is used to query a similarity search on Azure Cosmos DB. We want to take the 10 documents that are the most relevant to the user's query. This context is then given again to Google Vertex AI (with another model `gemini-2.5-flash-lite`). The final textual answer from Gemini 2.5 is sent back in the streamlit chat.



Based on the code provided for the previous lab, we refactored the scripts for the new infrastructure, keeping mostly the same scripts separation strategy. The chatbot has the same strategy as the previous one but is using the new services. For ease of execution, we have chosen a few hard-coded values inside `config.ini` to avoid choosing and giving names. All values that cannot be chosen in advance (such as the public IP or the folder where you want to download PDFs), are provided via CLI flags.

Prerequisites

This assumes you have Python installed and you have access to the 3 following clouds: Google Cloud, Microsoft Azure and Switch Engines.

1. Install Python dependencies

```
pip install -r requirements.txt
```

2. For Azure: [Install the Azure CLI](#) `az` and run `az login`.

3. Setup the Google Cloud environment

- i. Create manually a [New project on GCloud](#) named `chatbot` and get it's ID. For this example, we got `chatbot-475420`.
- ii. Change your `config.ini` file with your ID under key `vertexai` > `project_id`
- iii. Enable Google Vertex AI API [here](#)
- iv.

You have to create a Service account, go in IAM, and create one like on this picture

The screenshot shows the Google Cloud IAM & Admin console. The left sidebar contains a navigation menu with options like IAM, PAM, Principal Access Boundary, Organizations, Identity & Organization, Policy Troubleshooter, Policy Analyzer, Organization Policies, Service Accounts (highlighted), Workload Identity Federat..., Workforce Identity Federa..., Labels, Tags, Settings, Privacy & Security, Identity-Aware Proxy, and Roles. The main content area is titled 'Create service account' and shows a three-step wizard. Step 1, 'Create service account', includes fields for 'Service account name', 'Display name for this service account', 'Service account ID *' (with the value 'service-account'), and 'Email address: service-account@chatbot-475420.iam.gserviceaccount.com'. There is also a 'Service account description' field. A 'Create and continue' button is at the bottom of step 1. Steps 2 and 3 are 'Permissions (optional)' and 'Principals with access (optional)' respectively. At the bottom of the wizard are 'Done' and 'Cancel' buttons.

v. Add **Vertex AI User** role and Continue with defaults.

✓ Create service account

2 Permissions (optional)

Grant this service account access to chatbot so that it has permission to complete specific actions on the resources in your project. [Learn more](#)

Role
 Vertex AI User
 Grants access to use all resource in Vertex AI

IAM condition (optional) ?
 + Add IAM condition

+ Add another role

Help me choose roles

Continue

3 Principals with access (optional)

Done Cancel

i. Now you have a service account, click on it

Google Cloud chatbot Search (/) for resources, docs, products, and more Search

IAM & Admin / Service accounts

Service accounts + Create service account Delete

Service accounts for project "chatbot"

A service account represents a Google Cloud service identity, such as code running on Compute Engine VMs, App Engine apps, or systems running outside Google. [Learn more about service accounts.](#)

Organization policies can be used to secure service accounts and block risky service account features, such as automatic IAM Grants, key creation/upload, or the creation of service accounts entirely. [Learn more about service account organization policies.](#)

Filter Enter property name or value

Email	Status	Name	Actions
<input type="checkbox"/> service-account@chatbot-475420.iam.gserviceaccount.com	Enabled		

i. Go in the Keys tab to **Add key** > Create new key

service-account@chatbot-475420.iam.gserviceaccount.com

Service account list page Permissions **Keys** Metrics Logs Principals with access

Keys

Service account keys could pose a security risk if compromised. We recommend you avoid downloading service account keys. [Learn more about the best way to authenticate service accounts on Google Cloud.](#)

Google automatically disables service account keys detected in public repositories. You can customize this behavior using the 'iam.serviceAccountKeyExposureResponse' organization policy. [Learn more](#)

Add a new key pair or upload a public key certificate from an existing key pair.

Block service account key creation using [organization policies](#). [Learn more about setting organization policies for service accounts](#)

Add key

Type	Status	Key	Creation date	Expiration date
No rows to display				

i. Download as JSON format

Create private key for "service-account@chatbot-475420.iam.gserviceaccount.com"

Downloads a file that contains the private key. Store the file securely because this key can't be recovered if lost.

Key type

☒ JSON
Recommended

☐ P12
For backward compatibility with code using the P12 format

Cancel Create

i. Place this JSON file under a file named `vertexai-service-account-key.json` in the current folder

4. Setup the Switch Engines environment

i. [Go login on Switch Engines panel](#) to get your generated API password

ii. Make sure you have the LS (Lausanne) region chosen not the ZH one...

iii. Go into the API access page to get your `clouds.yaml`

Project / API Access

API Access

Displaying 14 items

Service	Service Endpoint
Cloudformation	https://cloudformation.api.unil.cloud.switch.ch/v1
Compute	https://nova.api.unil.cloud.switch.ch/v2.1
Dns	https://designate.api.zhdk.cloud.switch.ch
EC2	https://service-is-disabled

View Credentials Download OpenStack RC File

- OpenStack clouds.yaml File
- OpenStack RC File
- EC2 Credentials

iv. Move this file under a `switch` subfolder in this directory

v. Edit it to add your password

vi. An SSH keypair will be created automatically and downloaded under `switch/switchengine-tsm-cloudsys.pem` if it doesn't already exist.

vii. Change your `default` security group to accept port 22 (SSH) and 8501 (the chatbot).

Project / Network / Security Groups / Manage Security Group Rules

Manage Security Group Rules: default (79c2f4f1-2546-4893-b4af-87ba56e88960)

Displaying 6 items

Direction	Ether Type	IP Protocol	Port Range	Remote IP Prefix	Remote Security Group	Description	Actions
Egress	IPv4	Any	Any	0.0.0.0/0	-	-	Delete Rule
Egress	IPv6	Any	Any	::/0	-	-	Delete Rule
Ingress	IPv4	Any	Any	-	default	-	Delete Rule
Ingress	IPv4	TCP	22 (SSH)	0.0.0.0/0	-	-	Delete Rule
Ingress	IPv4	TCP	8501	0.0.0.0/0	-	Chatbot port	Delete Rule
Ingress	IPv6	Any	Any	-	default	-	Delete Rule

Displaying 6 items

S3-compatible container creation on Switch engines

We know this wasn't 100% recommended by John White but we tried and it works.

```
python manage-S3_switch.py --pdf_path ../../../../TSM_CloudSys-2024-25.pdf
```

With this script, we create container in object store, upload an pdf. We can also download this pdf, list object storage and contents and delete a dedicated container.

Vector database creation on Azure Cosmos DB service

We are going to use **Azure Cosmos DB** for this part. The home page of this service is here: [Create an Azure Cosmos DB account](#).

Note: the account name is also used as the database name. This name must be unique and can only

1. Activate the Cosmos DB service for your Azure account

```
az provider register --namespace Microsoft.DocumentDB
```

It may take 1–2 minutes to complete. You can check the registration status with the command:

```
az provider show --namespace Microsoft.DocumentDB --query "registrationState"
```

The command should return **Registered**

2. Then you can run

```
> python setup-azure.py
Provisioned resource group groupd-chatbot-deploy
Created Cosmos DB account: groupdchatbotd1234
```

3. Get the primary key in your portal and save it under a `azure-db-key.txt` file in this folder. This will be used for authentication for next scripts.

[Home](#) > [groupdchatbotd1234](#)

The screenshot shows the Azure portal interface for the Cosmos DB account 'groupdchatbotd1234'. The 'Keys' tab is selected, displaying the account's URI and a list of keys. The 'PRIMARY KEY' is highlighted with a red circle, and its value is masked with asterisks. The 'PRIMARY CONNECTION STRING' is also visible, masked with asterisks. The 'Read-only Keys' section is also present, showing a 'SECONDARY KEY' and its connection string. The left sidebar shows the account's settings, including 'Account Throughput', 'Features', 'Replicate data globally', 'Default consistency', 'Backup & Restore', 'Networking', 'CORS', 'Data Encryption', 'Dedicated Gateway', 'Keys', 'Advisor Recommendations', and 'Microsoft Defender for Cloud'.

Vectorizing the PDF Files

We want to download the files in S3 again, ask Google Vertex AI to generate embeddings and store them in a container in a database in Cosmos DB.

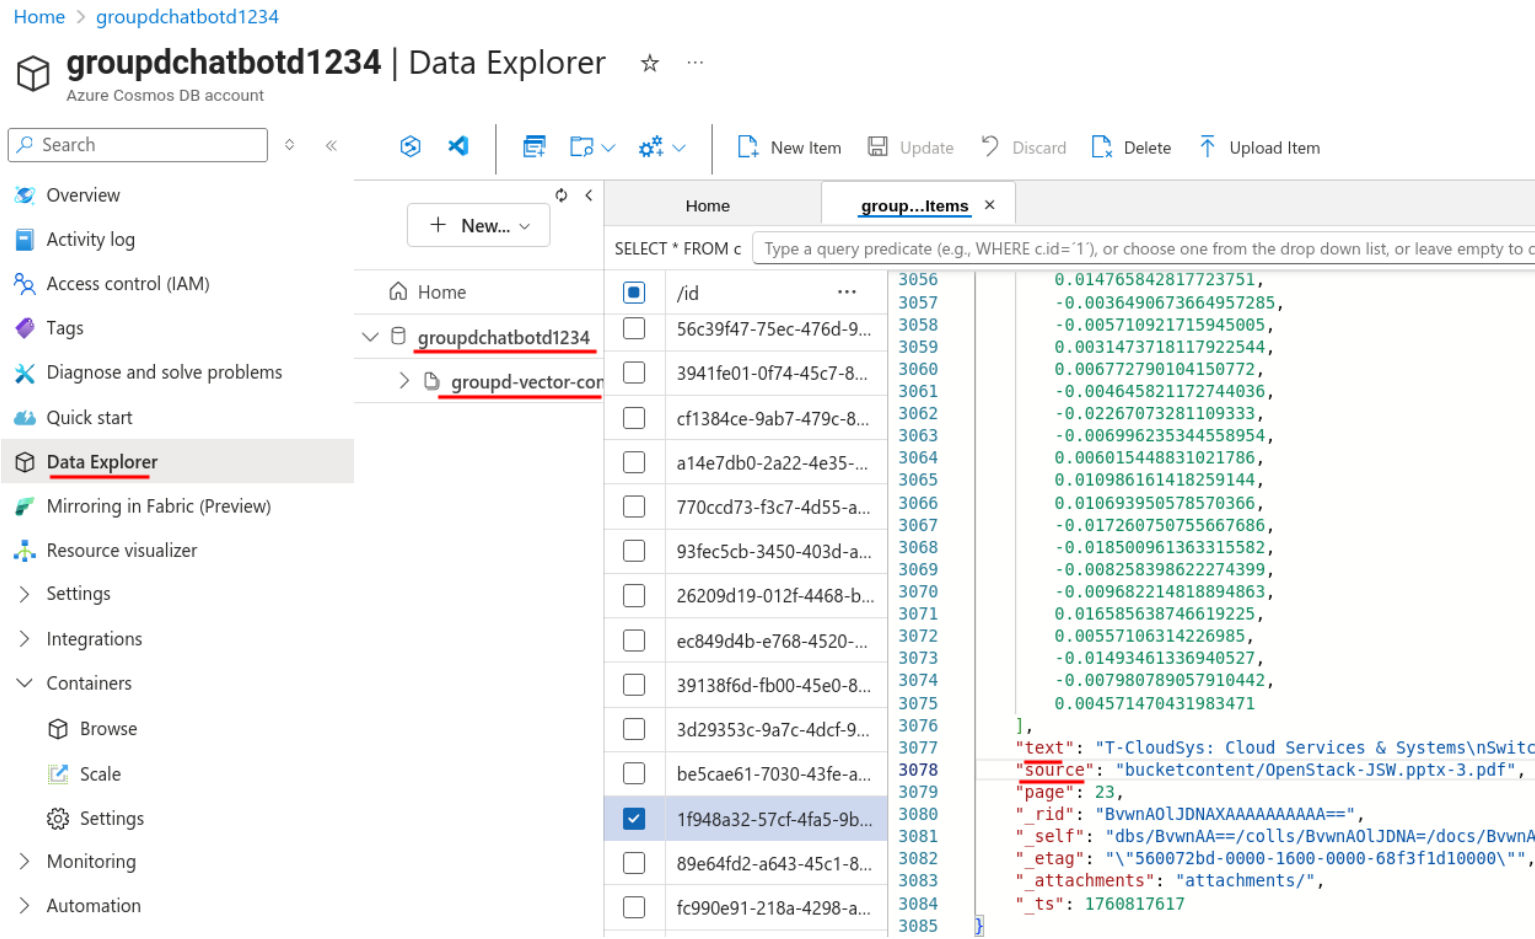
First download the pdf file from S3:

```
$ python manage-S3_switch.py --download
Download TSM_CloudSys-2024-25.pdf in ./s3-download/TSM_CloudSys-2024-25.pdf...
TSM_CloudSys-2024-25.pdf downloaded successfully.
```

And run the script `vectorise-store.py` which is an adaptation of the provided script in the previous lab.

```
python vectorise-store.py --local_path s3-download
```

You can verify the embeddings have been stored by looking at the Azure portal. You should see entries with chunks of the text, a giant float array, the text of the chunk and the source file.



Accessing the application locally

To make sure everything is working, there is a little `test.py` script that allows to quickly test the bot without open the streamlit chat. The message sent is hard-coded and can be changed if needed.

```
python test.py
```

If the AI is giving an answer and no Python stacktrace, it means everything is fully working. It's not possible to run the chatbot locally.

```
streamlit run chatbot.py
```

which should open your browser with a fully working bot. You can already ask questions about your vectorized PDFs.

Create the instance on Switch Engines

Make sure to fill the `config.ini` file !
These files will be used in the `manage_instance_switch.py` script, which you can simply run with

```
$ python manage_instance_switch.py --create
List Servers:
Create Server:

VM 'groupd-labo1' created.
You can login with SSH in a minute with
ssh -i ./switch/switchengine-tsm-cloudsys.pem ubuntu@86.119.31.138
```

Now an Ubuntu 22.04 instance of flavour `m1.small` should be running and accessible. It take 1-2 minutes to be accessible via SSH.

Project

API Access

Compute

Overview

Instances

Images

Key Pairs

Server Groups

Volumes

Project / Compute / Instances

Instances

Displaying 1 Item

	Instance Name	Image Name	IP Address	Flavor
<input type="checkbox"/>	groupd-labo1	Ubuntu Jammy 22.04 (SWI TCHengines)	10.0.5.49, 2001:620:5ca1:1f0:f816:3eff:fe69:ce01	m1.small

Instance ID =

Deploy the chatbot on the VM

The VM is ready to host the chatbot. Get the hostname showed previously and run it like that


```
$ python deploy.py --host 86.119.31.138
```

```
Uploading current folder into the VM
```

```
Starting deploy.sh on the VM
```

```
Hit:1 http://ch.archive.ubuntu.com/ubuntu jammy InRelease
```

```
Get:2 http://ch.archive.ubuntu.com/ubuntu jammy-updates InRelease [128 kB]
```

```
...
```

```
Defaulting to user installation because normal site-packages is not writeable
```

```
Collecting openstacksdk
```

```
  Downloading openstacksdk-4.7.1-py3-none-any.whl (1.8 MB)
```

```
1.8/1.8 MB 10.3 MB/s eta 0:00:00
```

```
Collecting azure-cosmos
```

```
  Downloading azure_cosmos-4.14.0-py3-none-any.whl (385 kB)
```

```
...
```

```
CHATBOT HAS BEEN DEPLOYED...
```

```
Please open http://86.119.31.138:8501 in your Web browser !
```

It installs Python, Pip, install the Pip dependencies, and start the chatbot in background.

You can finally open the link printed on the last line.

Chat with your lecture



what is openstack ?

S OpenStack is a free and open-source cloud-computing software platform. It provides services for managing a cloud environment on-the-fly and consists of a group of interrelated projects that control processing, storage, networking resources, authentication, and more. It was initially designed to provide Infrastructure as a Service (IaaS) but now offers "high-level" services as well.

Ask a question for your knowledge base



Delete the infrastructure

At the end of the lab, when you need to delete the Azure infrastructure, run this script (this can take several minutes to delete the resource group)

```
> python delete-azure.py
Cosmos DB account 'groupdchatbotdb1234' deleted successfully.
Resource group 'groupd-chatbot-deploy' deleted successfully.
```

You can delete the Switch VM and associated resource, with this command. Make sure to give the correct VM name.

```
> python manage_instance_switch.py --delete-vm groupd-labo1
List Servers:
groupd-labo1 - ACTIVE - fcef1b69-83a8-4173-a399-8dcdb3758bb3
Server 'groupd-labo1' deleted
Floating IP 86.119.31.63 released
```

You can delete the S3 from Switch engine

```
> python manage-S3_switch.py --delete  
Deleted object: TSM_CloudSys-2024-25.pdf  
Container 'groupd' deleted successfully.
```

The Google Cloud project `chatbot` could be deleted manually.