

# Kubernetes lab

---

**Authors:** Olivia Manz and Samuel Roland in **Group D**

The whole project is available on this Github repository: <https://github.com/samuelroland/CloudSys-labs/tree/main/lab5>

## Introduction

---

We have just done the Kubernetes exercise where the forecasting was deployed to a “Kind” cluster.

For the lab, you will repeat the Kubernetes exercise but deploy all the services to a k3s cluster instead of a “Kind” Cluster.

We will reuse the files from the previous exercise and deploy them on AWS.

## Create AWS instances

---

Sign in at AWS EC2 Console and create 4 instances:

- 1x: 2 CPU, 4GB RAM instance, 50GB, named `GroupD-control-pane`
- 3x: 1CPU, 2GB RAM at most, 30GB, named `control-worker-1`, `control-worker-2`, `control-worker-3`

A security group `launch-wizard-114` is automatically created. We will use the same key pair `Gd-control-plane.pem` for all instances.

To allow k3s to communicate between the master and worker nodes during installation, add a rule in the security group to allow TCP on port `6443`. We can prepare the futur by adding:

- `UDP 8472` to allow the forecast service to push data to Redis (node-to-node communication)
- `TCP 32000` to allow access to the Grafana dashboard from outside the cluster. To minimize exposure, restrict the source to the security group itself for `UDP 8472` and `TCP 6443`.

## Install k3s on control-pane

---

Connect to the `GroupD-control-plane` instance via SSH. The IP address (or SSH command) is available in the AWS console:

```
ssh -i "Gd-control-plane.pem" ec2-user@ec2-54-204-109-236.compute-1.amazonaws.com
```

Install k3s:

```
curl -sfL https://get.k3s.io | sh -
```

Verify the installation

```
$ sudo k3s kubectl get nodes
NAME           STATUS   ROLES          AGE    VERSION
ip-172-31-29-6.ec2.internal   Ready    control-plane, master   70s    v1.33.5+k3s1
```

Display the node token for adding workers later:

```
sudo cat /var/lib/rancher/k3s/server/node-token
```

# Install k3s on workers

Connect to each worker via SSH and install k3s in agent mode, using the control plane IP and token:

```
curl -sfL https://get.k3s.io | K3S_URL=https://<IP-control-plane>:6443 K3S_TOKEN=<token-node> sh -
```

Repeat for all workers. On the control plane, verify the nodes:

```
$ sudo k3s kubectl get nodes -o wide
NAME           STATUS   ROLES      AGE    VERSION   INTERNAL-IP   EXTERNAL-IP   OS-
IMAGE          KERNEL-VERSION
ip-172-31-23-123.ec2.internal   Ready     <none>    3m17s   v1.33.5+k3s1   172.31.23.123   <none>
Amazon Linux 2023.9.20251027   6.1.156-177.286.amzn2023.x86_64
ip-172-31-24-250.ec2.internal   Ready     <none>    2m22s   v1.33.5+k3s1   172.31.24.250   <none>
Amazon Linux 2023.9.20251027   6.1.156-177.286.amzn2023.x86_64
ip-172-31-27-254.ec2.internal   Ready     <none>    2m46s   v1.33.5+k3s1   172.31.27.254   <none>
Amazon Linux 2023.9.20251027   6.1.156-177.286.amzn2023.x86_64
ip-172-31-29-6.ec2.internal     Ready     control-plane,master 7m57s   v1.33.5+k3s1   172.31.29.6     <none>
Amazon Linux 2023.9.20251027   6.1.156-177.286.amzn2023.x86_64
containerd://2.1.4-k3s1
```

## Prepare nodes

As we should install grafana, redis and data-retrieval on the master and forecast on the workers, we label each node:

```
sudo k3s kubectl label node ip-172-31-29-6.ec2.internal role=control-plane
sudo k3s kubectl label node ip-172-31-23-123.ec2.internal role=worker
sudo k3s kubectl label node ip-172-31-24-250.ec2.internal role=worker
sudo k3s kubectl label node ip-172-31-27-254.ec2.internal role=worker
```

We can check the result:

```
$ sudo k3s kubectl get nodes --show-labels
NAME           STATUS   ROLES      AGE    VERSION   LABELS
ip-172-31-23-123.ec2.internal   Ready     <none>    3m28s   v1.33.5+k3s1   beta.kubernetes.io/
arch=amd64,beta.kubernetes.io/instance-type=k3s,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/
hostname=ip-172-31-23-123.ec2.internal,kubernetes.io/os=linux,node.kubernetes.io/instance-type=k3s,role=worker
ip-172-31-24-250.ec2.internal   Ready     <none>    2m33s   v1.33.5+k3s1   beta.kubernetes.io/
arch=amd64,beta.kubernetes.io/instance-type=k3s,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/
hostname=ip-172-31-24-250.ec2.internal,kubernetes.io/os=linux,node.kubernetes.io/instance-type=k3s,role=worker
ip-172-31-27-254.ec2.internal   Ready     <none>    2m57s   v1.33.5+k3s1   beta.kubernetes.io/
arch=amd64,beta.kubernetes.io/instance-type=k3s,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/
hostname=ip-172-31-27-254.ec2.internal,kubernetes.io/os=linux,node.kubernetes.io/instance-type=k3s,role=worker
ip-172-31-29-6.ec2.internal     Ready     control-plane,master 8m8s   v1.33.5+k3s1   beta.kubernetes.io/
arch=amd64,beta.kubernetes.io/instance-type=k3s,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/
hostname=ip-172-31-29-6.ec2.internal,kubernetes.io/os=linux,node-role.kubernetes.io/control-plane=true,node-
role.kubernetes.io/master=true,node.kubernetes.io/instance-type=k3s,role=control-plane
```

As asked in the instructions, we want to have all the pods on the control-plane node and except for the Forecast pods on the worker nodes. To achieve this, we changed the deployment yaml files and to add a `nodeSelector` that makes it possible to select the node by role.

```
spec:
  template:
    spec:
      nodeSelector:
        role: control-plane
```

For `forecast-deployment.yaml`, we did the same thing with `role: worker`.

For Grafana, we changed the service type to `NodePort` to allow external access:

```
apiVersion: v1
```

```
kind: Service
metadata:
  name: grafana-service
spec:
  type: NodePort
  ports:
  - port: 3000
    targetPort: 3000
    nodePort: 32000 # port available from outside
  selector:
    app: grafana-pod
```

## Deploying the Services

We decide to copy deployment files on control-plane node and run these commands locally to upload the folder `deployment` under `/home/ec2-user/`.

```
scp -i Gd-control-plane.pem -r deployment ec2-user@54.204.109.236:/home/ec2-user/
```

On the control-plane node, we deployed Redis:

```
sudo k3s kubectl apply -f /home/ec2-user/deployment/redis-deployment.yaml
```

Like before, we created a secret for AWS credentials

```
kubectl create secret generic <name-of-the-secrets> \
--from-literal=AWS_ACCESS_KEY_ID=<Your access key> \
--from-literal=AWS_SECRET_ACCESS_KEY=<your secret key>
```

We then deployed Data-Retrieval:

```
sudo k3s kubectl apply -f /home/ec2-user/deployment/data-retrieval-deployment.yaml
```

We check that data-retrieval has finished its execution before deploying the forecast module.

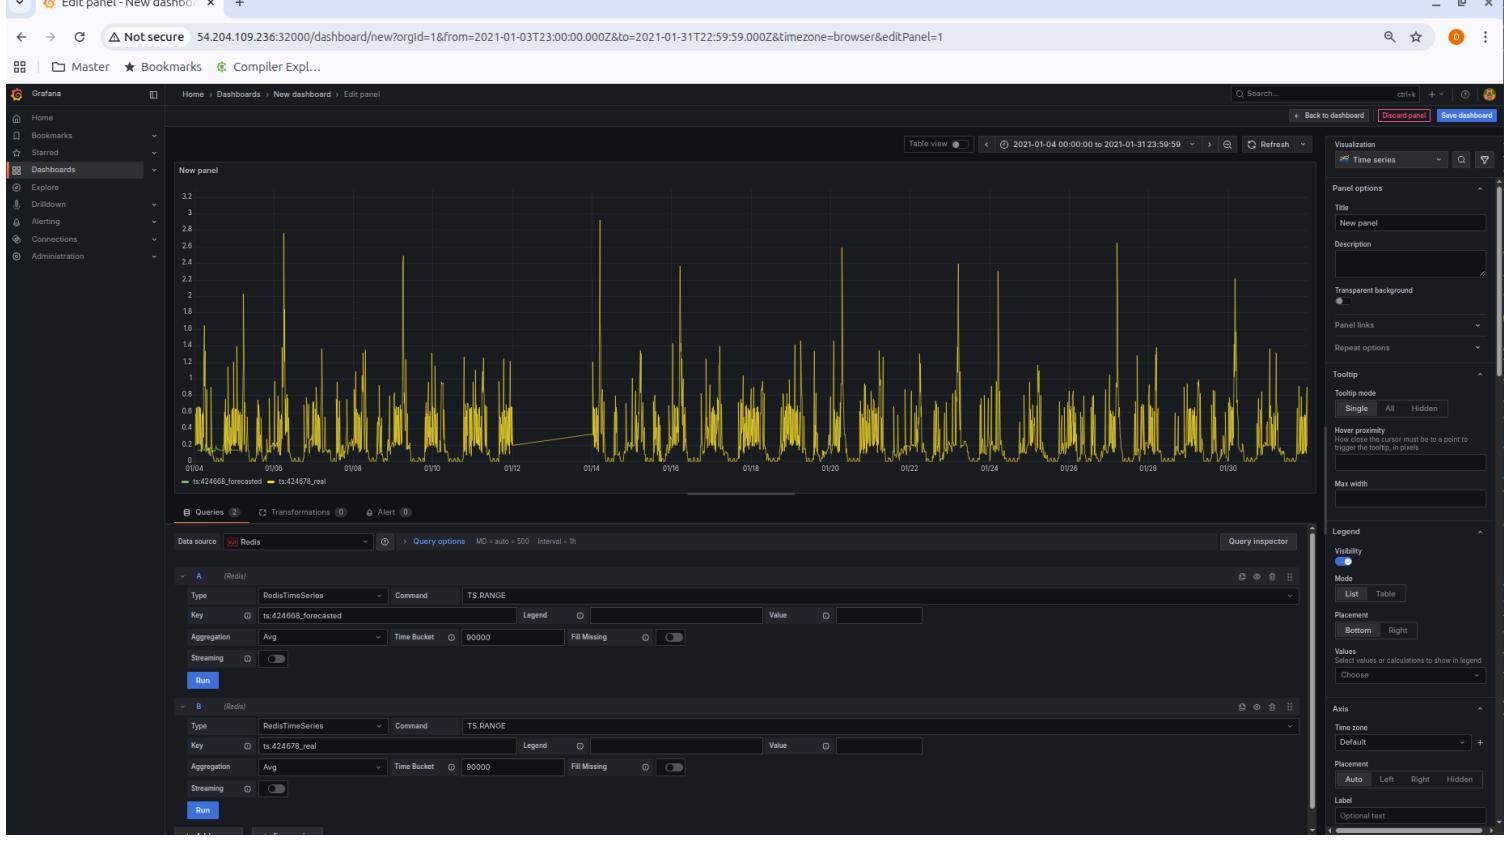
```
sudo k3s kubectl apply -f /home/ec2-user/deployment/forecast-deployment.yaml
```

Finally, we deployed Grafana:

```
sudo k3s kubectl apply -f /home/ec2-user/deployment/grafana-deployment.yaml
```

We verified the status of all pods and here is the full output of `sudo kubectl get all -o wide` in screenshot.

We accessed Grafana through the public IP of the control-plane node on port 32000. The following screenshot shows the result after the dashboard configuration.



## Difficulty

---

The main difficulty during this lab was knowing whether a rule was necessary to add to the security group. Particularly for the rule allowing the Forecast service to send data to Redis.